

Project Euler
- Prematurely optimized and Overengineered

Øistein Søvik

June 4, 2017

Contents

Project Euler 1: Multiples of 3 and 5	3
Project Euler 4: Largest palindrome product	6
Project Euler 5: Smallest multiple	7
Project Euler 7: 10 001st prime	8
Project Euler 8: Largest prime factor	9
Project Euler 3: Largest prime factor	10
Project Euler 10: Summation of primes	11
Project Euler 13: Large Sum	12

Project Euler 1: Multiples of 3 and 5

If we list all the natural numbers below 10 that are multiples of 3 or 5, we get 3, 5, 6 and 9. The sum of these multiples is 23.

Find the sum of all the multiples of 3 or 5 below 1000.

It is straight forward to iterate over the first 1000 numbers and check for divisibility.

```
def divisible_by_3_or_5(limit = 1000):
    count = 0
    for num in xrange(1, limit):
        if num%3 == 0 or num%5 == 0:
            count += num
    return count
```

The code above takes on average 12 μ s to run. Well under the arbitrary 1 s rule. However we can do better, a small improvement is to allow for different numbers than 3 or 5 and also allow for a customizable range.

```
def numbers_divisible(divisors=[3, 5], start=0, stop=100):
    count = 0
    for num in xrange(start, stop):
        for d in divisors:
            if num % d == 0:
                count += num
                break
    return count
```

It could seem the above code is quite fast however running `numbers_divisible([3, 5], 0, 10**9)` takes about 450 s to run. This is painfully slow, but also expected since the code runs in $\mathcal{O}(n)$. As we shall see we can make the code run in constant time.

Improved algorithm for two numbers The first step is to count the number of numbers divisible by 3, then 5. However adding these two will give the incorrect answer, we can see why by listing the first few numbers. [3, 6, 9, 15] and [5, 10, 15]. So every number which is divisible by 3 and 5 is counted twice. Figure 1 visualizes this.




Figure 1

```
div_by_3 = [3*i for i in range(int(999/3))]
div_by_5 = [5*i for i in range(int(999/5))]
div_by_3_and_5 = [15*i for i in range(int(999/15))]
return div_by_3 + div_by_5 - div_by_3_and_5
```

This is slightly faster than the naive implementation. However we can find a closed formula for the number of numbers divisible by 3, or any other number. The closed formula for the first n numbers are¹

$$0 + 1 + \dots + (m - 1) + m = \frac{m(m + 1)}{2}$$

To find the sum of the natural numbers starting at some number n up to some number m we can subtract the sum of numbers up to $n - 1$.

$$n + (n + 1) + \dots + (m - 1) + m = \frac{m(m + 1)}{2} - \frac{(n - 1)n}{2} = \frac{1}{2}(m + n)(m - n + 1)$$

Since $3 + 6 + 9 + 12 + \dots = 3(1 + 2 + 3 + \dots)$ we can just multiply the above formula by 3 to sum the multiples of 3. However we have to make sure we are not multiplying numbers larger than 1000. So we can not let m be 1000 and sum up a thousand multiples of 3. The largest of these would be $3 \cdot 1000$ which is 3 times as large as the limit. One solution is to take $m = 999/3$.

If we are summing over multiples of k , then the upper limit would be $(\text{textlimit} - 1)/k$ rounded down.

One problem with this code can be seen with the [6, 9]. The code removes multiples of $6 \cdot 9 = 56$ however the first value that is counted twice is 18. One way to solve this is to divide the product by the lowest common divisor (gcd). So we would have $56/\text{gcd}(6, 9) = 56/3 = 18$.

```
from fractions import gcd
def sum_divisible_by_k(k, start, limit):
    stop = int((limit-1)/float(k))
    return k*(stop+start)*(stop-start+1)/2

def divisible_by_a_or_b(num_a, num_b, start=0, limit=100):
    divisors = [num_a, num_b]
    total = 0
    for divisor in divisors:
        total += sum_divisible_by_k(divisor, start, limit)
    product = divisors[0]*divisors[1]/gcd(a, b)
    return total - sum_divisible_by_k(product, start, limit)
```

For values up to 10^5 the improved version runs around 4500 times faster. However as stated this function runs in constant time so a speed comparison is not really necessary. We have now done the *silly* speed improvements.

More than two numbers There are some bumps to work out for the generalized version. We shall start looking into the simplest case [2, 3, 5]. We can visualize the double counting as follows As we can see from fig. 2 we have to remove the multiples of $2 \cdot 3$, $2 \cdot 5$ and $3 \cdot 5$. However removing these will remove too many values. As an exercise you can check that 60 is missing if we do not add the multiples of $2 \cdot 3 \cdot 5$.

The next problem to work out is multiples. Take [2, 3, 8], there is no need to add 8 once we have added 2. The following code removes all multiples, all doubles and sorts the divisors.

```
def remove_multiples(divisors):
    new_divisors = []
    divisors = sorted(set(divisors))
    for divisor in divisors:
        if not any(divisor % d == 0 for d in new_divisors):
            new_divisors.append(divisor)
    return new_divisors
```

The final problem is that we still have to divide by the greatest common divisor. However we now have to perform this on a list of numbers. Luckily the gcd functions is additive so that $\text{gcd}(a, b, c) = \text{gcd}(\text{gcd}(a, b), c)$. Since we this be done quite a few times I will use memoization and save some calculation.

¹See https://en.wikipedia.org/wiki/1_%2B_2_%2B_3_%2B_4_%2B_%E2%8B%AF for further details.




Figure 2

```
gcd_dict = defaultdict(int)
def gcd_list(numbers):
    numbers_len = len(numbers)
    a = perm[0]
    for i in range(1, numbers_len):
        b = perm[i]
        if gcd_dict[(a, b)] == 0:
            gcd_dict[(a, b)] = gcd(a, b)
        a = gcd_dict[(a, b)]
    return a
```

The final code in all it's glory can now be written as

```
def sum_divisible_by_numbers(numbers, start, stop):
    divisors = remove_multiples(numbers)

    total = 0
    for divisor in divisors:
        total += sum_divisible_by_k(divisor, start, stop)

    k = -1
    for i in range(2, len(divisors)+1):
        for perm in combinations(divisors, i):
            product = reduce(mul, perm)/gcd_list(perm)
            total += k*sum_divisible_by_k(product, start, stop)
        k *= -1
    return int(total)
```

The running time of this algorithm is $\mathcal{O}(2^d)$ where d is the number of unique numbers to check for. This is also the reason we take time to remove duplicates and multiples of values.

Project Euler 2: Even Fibonacci numbers

Each new term in the Fibonacci sequence is generated by adding the previous two terms. By starting with 1 and 2, the first 10 terms will be:

1, 2, 3, 5, 8, 13, 21, 34, 55, 89, ...

By considering the terms in the Fibonacci sequence whose values do not exceed four million, find the sum of the even-valued terms.

```
sum(int(.5+5**(-.5)*(.5+.5*5**.5)**(3*n)) for n in
range(100) if int(.5+5**(-.5)*(.5+.5*5**.5)**(3*n)) < 4*10**6)
```

This is one of my favourite problems and as a teaser one of the solution can be written as

```
def sum_even_fibonacci(limit):
    a, b = 0, 2
    while b < limit:
        a, b = b, 4 * b + a
    return (a + b - 2) / 4
```

However to understand how this code works we have to go pretty far down the rabbit hole. A warning before we start: this problem is very easy to solve under 1s, and any improvements beyond this is purely for the amusement of the Author. By the end we will have developed a method to find the sum of all even Fibonacci numbers under 10^{10000} in about a 1ms.

Definition 1. *the Fibonacci numbers are the numbers in the following integer sequence, called the Fibonacci sequence:*

0, 1, 1, 2, 3, 5, 8, 13, 21, 34, 55, 89, 144, ...

The numbers have in common that each subsequent number is the sum of the previous two. In mathematical terms we can define the Fibonacci numbers as:

$$F_n = F_{n-1} + F_{n-2} \tag{1}$$

with initial conditions $F_0 = 0$ and $F_1 = 1$.

Note that this definition differs slightly from the proposed by Project Euler. The difference is that our indices is shifted $n \rightarrow n + 2$ such that $F_{\text{PEN}} = F_{n+2}$. The only reason for this is that it leads to slightly nicer notation later on.

The definition above leads us straight into the following code

```
def fib(n):
    if n < 2:
        return n
    else:
        return fib(n-1) + fib(n-2)

def sum_even_fib_naive(limit):
    n = total = 0
    while fib(n) < limit:
        n += 1
        if fib(n) % 2 == 0:
            total += fib(n)
    return total
```

This is so far the slowest solution I have yet written to a Project Euler problem. To just see how slow it is we can measure the number of times `fib` is called. To evaluate F_5 we have to call our function 35 times! See ?? for a visualization of this. If you have a bit of sparetime try to see how big this tree gets for F_6 , and if you have a couple of lifetimes to spare you can verify that `f(33)` makes 29 860 703 function calls. This




Figure 3: Shows the computations needed to evaluate F_5 .

is the reason our code is so slow, and it can be shown that `fib(n)` grows as $\mathcal{O}(\phi^n)$. Here ϕ is the golden ratio and shortly see how this number plays a role in our calculations. One method to fix the problem with recursively calling our function is to store the previously calculated values in memory

<pre>def fib_memo(n): pad = {0:0, 1:1} def func(n): if n not in pad: pad[n] = func(n-1) + func(n-2) return pad[n] return func(n)</pre>	<pre>from functools import lru_cache @lru_cache(maxsize=None) def fib(n): if n < 2: return n return fib(n-1) + fib(n-2)</pre>
--	---

Note: If you are using Python 2.7 you need to use `from functools32 import lru_cache` to use `@lru_cache`. Both of these functions speeds up `sum_even_fib_naive` dramatically. However the drawback of using memoization is that it needs to store a large number of values and thus is not very memory efficient. In this problem however memory is not a concern since to obtain the solution we only need to store 32 values in memory. On the other hand it is a good practice not to hog more memory than needed and in this problem it is easy to do so. ?? is know as a *second order linear homogeneous recurrence relation*. That was a bit of a mouthful, but it basically means we can find a closed form for the n th Fibonacci number. One way to find the solution to these equations is to simply *guess*. You are free to try $f(n) = An + B$, $f(n) = An^2 + Bn + C$ or any other polynomial. However none of these attempts work out. For recurrence relations our first guess is usually on the form $f(n) = A \cdot P^n$, inserting this in ?? gives

$$A \cdot (P^n - P^{n-1} - P^{n-2}) = 0 \quad (2)$$

For the relation to hold the left-hand side to be zero for all values of n either A or $P^n - P^{n-1} - P^{n-2}$ has to be zero. $f(n) = 0$ satisfies $F_n = F_{n-1} - F_{n-2}$, however it is obviously not the solution we are looking for. Hence we assume that $A \neq 0$, and we can safely divide ?? by A . In the same vein we also multiply the equation by P^{2-n}

$$P^2 - P - 1 = 0$$

This is called the *characteristic polynomial* to our recurrence relation. The roots of this equation gives us the valid values for P such that $f(n) = A \cdot P^n$ satisfies ??.

$$P = 0 \vee P = \frac{1 \pm \sqrt{5}}{2}$$

Again $P = 0$ is the *trivial* solution. I will leave it to you to verify that $f(n) = A[(1 + \sqrt{5})/2]^n$ satisfies $f(n) = f(n-1) + f(n-2)$. There is a small problem with our solution however, even though it satisfies the recurrence relation the *initial-values* are wrong. We need $f(0) = 0$, but then $f(0) = A$, hence $A = 0$ and we are again left with the trivial solution.

For a moment assume that both $f(n) = AP^n$ and $g(n) = BQ^n$ satisfies $??$, then $h(n) = AP^n + BQ^n$ will also satisfy the equation. Said more plainly *a sum of solutions is also a solution*. Our next guess therefore becomes

$$f(n) = A \left(\frac{1 + \sqrt{5}}{2} \right)^n + B \left(\frac{1 - \sqrt{5}}{2} \right)^n$$

What remains is finding A , and B such that $f(0) = 0$ and $f(1) = 1$. This leads to the following system of equations

$$A + B = 0 \quad \text{and} \quad A \left(\frac{1 + \sqrt{5}}{2} \right) + B \left(\frac{1 - \sqrt{5}}{2} \right) = 0$$

Solving this system is quite trivial and yields $A = 1/\sqrt{5}$ and $B = -A = -1/\sqrt{5}$.

$$F_n = \frac{1}{\sqrt{5}} \left(\frac{1 + \sqrt{5}}{2} \right)^n - \frac{1}{\sqrt{5}} \left(\frac{1 - \sqrt{5}}{2} \right)^n$$

This equation can be written in a variety of ways and I have collected a handful below

$$F_n = \frac{\phi^n - \psi^n}{\phi - \psi} = \frac{\phi^n - \psi^n}{\sqrt{5}} = \frac{\phi^n - (-\phi)^{-n}}{\sqrt{5}} \quad (3)$$

The notation $\phi = (1 + \sqrt{5})/2$ and $\psi = (1 - \sqrt{5})/2$ was used for convenience sake². Proving the last relation implies proving $\psi = -1/\phi \implies \psi^n = (-1/\phi)^n = \phi^{-n}$ is left as a small exercise. The implications this has is that $(-\phi)^{-n}$ quickly diminishes $(-\phi)^{-10} \approx 0,00813$. Hence we can implement the F_n function as $F_n = \text{ceil}(\phi^n/\sqrt{5})$.

```
SQRT_5 = 5**0.5
PHI = 0.5*(1 + SQRT_5)
def fib_exact(n):
    return int(0.5 + (PHI**n)/SQRT_5)
```

Where `int(0.5+num)` was used to round up, since $\text{floor}(x + 0.5) = \text{ceil}(x)$. At a first glance this seems like the perfect solution as it runs in constant running time and for our problem it works without quirks. However exponentiation with decimal numbers is very prone to rounding errors. Checking the accuracy of this function one can see that it starts deviates from the actual Fibonacci numbers after $n = 71$. For now I present an alternative. The following code is a simple exact way to generate the fibonacci numbers without relying on recursion nor memoization.

```
def fib_generator():
    F0, F1 = 0, 1
    while True:
        yield F0
        F1, F0 = F1 + F0, F1

def sum_even_fibonacci(limit):
    total = 0
    for nth_fib in fib_generator():
        if nth_fib % 2 != 0: continue
        if nth_fib > limit: return total
        total += nth_fib
```

Even though the closed form could not be used our efforts have not been completely wasted. $??$ gives us an excellent starting point in figuring out how many terms we need. At first glance this seems complicated

²Know as the golden ratio

and for complete accuracy we should use a numerical method such as Newton-rhapson or the bisection method. However since $(-\phi)^{-n}$ diminishes so quickly a rough solution is just to ignore it.

$$M = F_{3n} \Rightarrow M \approx \frac{\phi^{3n}}{\sqrt{5}} \Rightarrow n \approx \frac{1}{6} \log_{\phi}(5) + \frac{1}{3} \log_{\phi}(M)$$

Where M denotes the largest F_n we will allow. For every $M > 1$ the approximation is equal to the actual value rounded down.

```
from math import log

PHI = (1 + 5**0.5)/float(2)
LOG_5 = log(5, PHI)/float(6)

def largest_even_fib_index(n):
    return int(LOG_5 + log(n, PHI)/float(3))

def sum_even_fibonacci_w_end(M):
    F0, F1, total = 0, 1, 0
    for _ in range(largest_even_fib_index(M)):
        for _ in range(3):
            F1, F0 = F1 + F0, F1
            total += F0
    return total
```

Note that even after all these efforts the running time of our algorithm is still the same. By writing out the first Fibonacci numbers we can see a pattern start to form

0, 1, 1, 2, 3, 5, 8, 13, 21, 34, 55, 89, 144, ...

It seems that every third Fibonacci number is divisible by 3. We can write a short proof using induction to verify this suspicion. We want to prove that F_{3n} is divisible by 2 for every $k \geq 0$.

Proof. The base case is that F_0 is divisible by 2, which it is since $F_0 = 0$. Assume that there exists some k such that F_{3k} is divisible by 2. We have to prove that this implies that $F_{3(k+1)} = F_{3k+3}$ is divisible by 2.

$$F_{3+3k} = F_{2+3k} + F_{1+3k} = [F_{3k+1} + F_{3k}] + F_{3k+1} = F_{3k} + 2F_{3k+1} \quad (4)$$

Where ?? was used twice. Since F_{3k} is divisible by 2 from the induction hypothesis and $2F_{3+3k}$ is clearly divisible by 2, the claim follows by induction. \square

Since there is a closed form for the Fibonacci numbers, it is not unreasonable to think that there exists a similar pattern for the even Fibonacci numbers, but how can we discover such a pattern? ?? gives us a starting point, since it has a relation between two even Fibonacci numbers: $F_{3(k+1)}$ and F_{3k} . The only thing that remains is to rewrite $2F_{3k+1}$ as a sum of even Fibonacci numbers.

$$\begin{aligned} 2F_{3k+1} &= 2(F_{3k} + F_{3k-1}) \\ &= 2F_{3k} + 2F_{3k-1} \\ &= 2F_{3k} + F_{3k-1} + F_{3k-2} + F_{3k-3} \\ &= 3F_{3k} + F_{3(k-1)} \end{aligned}$$

Replacing $2F_{3k+1}$ with $3F_{3k} + F_{3(k-1)}$ in ?? gives us the following relation

$$F_{3(k+1)} = 4F_{3k} + F_{3(k-1)} \quad (5)$$

This can also be written as $E(n) = 4E(n-1) + E(n-2)$ when we let $E(n)$ denote the n th even Fibonacci number. Interestingly enough ?? also holds for the odd Fibonacci numbers, and one can prove that $F_n = 4F_{n-3} + F_{n-6}$ holds for every n . This however will be left as an exercise for the reader.

```
def sum_even_fast(limit):
    a, b = 0, 2
    total = 0
    while b < limit:
        total += b
        a, b = b, 4 * b + a
    return total
```

This code is much faster than our previous attempts, and could further be slightly optimized by using `for _ in range(largest_even_fib_index(M)):` instead of `while b < limit:`.

The next optimization is to find a closed form for the sum of the first n Fibonacci numbers.

Proposition 1.

$$\sum_{i=0}^n F_{3i} = \frac{F_{3n+2} - 1}{2} = \frac{F_{3n} + F_{3n+2} - 2}{4}$$

Proof. We will prove the first relation by induction, while the second is left as an exercise. Base case: $F_{3 \cdot 0} = 0$ and $(F_{0+2} - 1)/2 = 0$. Assume that for some k we have $\sum_{i=0}^k F_{3i} = \frac{F_{3k+2}-1}{2}$. We want to prove that this implies that $\sum_{i=0}^{k+1} F_{3i} = \frac{F_{3(k+1)+2}-1}{2}$ holds.

$$\begin{aligned} \sum_{i=0}^{k+1} F_{3i} &= F_{3(k+1)} + \sum_{i=0}^k F_{3i} = F_{3(k+1)} + \frac{F_{3k+2} - 1}{2} \leftarrow \text{used the induction hypothesis} \\ &= \frac{F_{3(k+1)} + F_{3k+3} + F_{3k+2} - 1}{2} = \frac{F_{3(k+1)} + F_{3(k+1)+1} - 1}{2} = \frac{F_{3(k+1)+2} - 1}{2} \end{aligned}$$

Since $F_{3k+3} + F_{3(k+1)}$ and $F_{3k+2} + F_{3k+3} = F_{3k+4} = F_{3(k+1)+1}$, this concludes the proof. \square

```
def constant_even_fib(limit):
    n = largest_even_fib_under_n(limit)
    return (int(0.5+PHI**((3*n+2)/5**0.5)-1))/2
```

This is a rather neat code, however it still suffers from floating point errors, when the powers get too high. Another solution is to find a closed formula for $E_n = 4E_{n-1} + E_{n-2}$. Similar as we did in equation nigga we can guess that $E_n = A \cdot P^n + B \cdot Q^n$, where P and Q are the roots of the characteristic polynomial $x^2 - 4x - 1 = (x - 2)^2 - 5$.

$$E_n = A(2 + \sqrt{5})^n + B(2 - \sqrt{5})^n$$

As before we can find the constants using $E_0 = 0$ and $E_1 = 2$.

$$\begin{aligned} E_n &= \frac{1}{\sqrt{5}}(2 + \sqrt{5})^n - \frac{1}{\sqrt{5}}(2 - \sqrt{5})^n \\ \frac{E_n + E_{n+1} - 2}{4} &= \frac{1}{4} \left(\frac{1}{\sqrt{5}}(2 + \sqrt{5})^n + \frac{1}{\sqrt{5}}(2 + \sqrt{5})^{n+1} - 2 \right) = \frac{1}{4\sqrt{5}}(3 + \sqrt{5})(2 + \sqrt{5})^n - \frac{1}{2} \end{aligned}$$

Since $|(2 - \sqrt{5})^n| < 1/2$ for all $n > 1$, we only need to add the powers of $2 + \sqrt{5}$.

```
def fib_sum_even_constant(limit):
    n = largest_even_fib_under_n(limit)
    return int(.5+5**(-.5)*((3+5**0.5)*(2 + 5**0.5)**n)/4
```

Which of these two algorithms if any do you think performs the best? Note that the -0.5 was dropped again because how Python performs rounding. This concludes how far we can go in constant time, however we are not quite at the end of the road.

2.1 Fibonacci in sublinear time

If we want to generate the first n Fibonacci numbers the fastest we can do is in polynomial time $\mathcal{O}(n)$. This is only natural since we have to iterate over every element. However as we will see if we are only interested in a particular F_n we do not need to generate every F_i where $0 \leq i \leq k$. There are many different ways to achieve this, I will briefly show three algorithms.

There have been many smaller improvements on the Lucas numbers and the matrix multiplication method. One is credited to ? which proposed the following algorithm.

```
def fib_takahashi(n):
    if n == 0:
        return n
    F, L, sign, exp = 1, 1, -2, int(log(n, 2))
    mask = 2**exp
    for i in xrange(exp - 1):
        mask = mask >> 1
        F2 = F**2
        FL2 = (F + L)**2
        F = ((FL2 - 6*F2) >> 1) - sign
        if n & mask:
            temp = (FL2 >> 2) + F2
            L = temp + (F << 1)
            F = temp
        else:
            L = 5*F2 + sign
            sign = -2 if n & mask else 2
    if n & (mask >> 1) == 0:
        return F * L
    else:
        return ((F + L) >> 1) * L - (sign >> 1)
```

I will not go into details on how the code works as this is explained much better than I can in his paper, see [?].

References

- [?] Daisuke Takahashi. A fast algorithm for computing large fibonacci numbers. *Information Processing Letters*, 75(6):243 – 246, 2000. ISSN 0020-0190. http://www.math.tamu.edu/~snpolloc/math491_691/Takahashi00.pdf.

Project Euler 3: Largest prime factor

The prime factors of 13195 are 5, 7, 13 and 29.

What is the largest prime factor of the number 600851475143?

This is another rabbit hole that goes deep, so be warned. However we will stop before you accidentally ends up with a masters degree in mathematics. While checking if a number is prime is easy, prime factorization on the other hand is hard. Said very naively this is because we do not need to find a single factor to ensure that a number is prime. Instead the most common technique is a probabilistic approach. However we will come back to that later.

Instead we will look at a few simple algorithms for factoring primes. At the end I will give some pointers on more complicated algorithms which is used to factor bigger and harder composite numbers.

Note that the problem only asks us to find the biggest factor. However in term of running time this is just as hard as finding all the factors. The most naive approach can be implemented as follows.

```
def largest_primefactor_naive(num):  
    i = 2  
    while i < num:  
        while num % i == 0:  
            num //= i  
        i += 1  
    return num
```

We can improve the running time in several ways. The first is an improvement of the upper bound

Lemma 1. *Every composite number n has exactly one prime factor q such that $q > \sqrt{n}$.*

Proof. We will prove this using contradiction. Since n is composite, n has *atleast* two prime factors p and q . Assume by contradiction that $p > \sqrt{n}$ and $q > \sqrt{n}$. This implies that $p \cdot q > \sqrt{n} \cdot \sqrt{n} > n$. We have now reached a contradiction because: p and q are prime factors of $n \iff n = p \cdot q$. \square

?? implies that we only need to test numbers up to \sqrt{n} . The last factor $q > \sqrt{n}$ must be a prime factor otherwise we could have written it as $p = a \cdot b$ and $a, b < \sqrt{n}$. So a, b would have been discovered as factors since we test numbers up to \sqrt{n} .

Another improvement is that we do not need to check *every* number up to \sqrt{n} . If we first check that the number is divisible by 2 we only need to iterate over the odd numbers. Which of these two methods do you think gives the greatest improvement in speed? From ?? we can see that the different wheel-

Table 1: Timings in ms for some of the prime factor methods

Number	Naive	sqrt(n)	wheel-2	wheel-23	wheel-235
600 851	4,954	0,141	0,031	0,030	0,012
60 085 147	124,834	1,246	0,214	0,237	0,117
6 008 514 751	1047,408	2,806	0,799	0,651	0,416
600 851 475 143	5,252	2,729	0,594	0,544	0,293
6 008 514 751 436 008	timeout	158,620	71,072	81,564	47,053

factoring algorithms outperforms the naive implementation by several orders of magnitude. It is difficult to see exactly how much of an improvement **wheel-23** is over **wheel-2**, because our numbers are too small. However **wheel-235** offers a significant improvement. We can generalize the **verb** algorithm to n -primes using the following code

```

from primesieve import generate_n_primes

def largest_factor(n, number_of_spokes = 3):

    spokes = generate_n_primes(number_of_spokes)
    wheel = reduce(lambda x, y: x * y, spokes)
    composites = []
    for i in range(1, wheel):
        if all( i % k != 0 for k in spokes):
            composites.append(i)

    for x in spokes:
        if n % x == 0 and n > x:
            n //= x
            while n % x == 0:
                n //= x
            if n == 1:
                return x

    x = 0
    limit = int(n**0.5) + 1
    while x < limit:
        for k in composites:
            if n % (x+k) == 0 and x+k > 1:
                n //= (x+k)
                while n % (x+k) == 0:
                    n //= (x+k)
                if n == 1: return x+k
                limit = int(n**0.5) + 1
        x += wheel
    return n

```

We can improve the wheel factorization further by only iterating over the primes. This is the best we can do for a wheel factorization, however it is only best given that we have a fast way to generate primes. Keeping a large list of primes would slow this method down significantly. Luckily Python has a few libraries which can help us solve this problem.

```

from primesieve import Iterator

def prime_gen(num):
    if isprime(num): return num
    it = Iterator()
    prime = it.next_prime()
    limit = num**0.5
    while prime < limit:
        if num % prime == 0:
            num //= prime
            while num % prime == 0: num //= prime
            if num == 1: return prime
            if isprime(num): return num
            limit = num**0.5
        prime = it.next_prime()
    return num

```

As we shall see for small numbers this gives a significant speedboost. Until now all of our methods have consisted of sieving out primes starting with the smallest one. As you probably have seen this is slow for all but the tiniest of numbers. To get another speed improvement we need to change our approach. One of the

simplest alternatives to the wheel-factorization is the Pollard's rho algorithm ³.

Pollard's rho algorithm

```
def pollard_rho(N):
    if N%2==0:
        return 2
    x = randint(1, N-1)
    y = x
    c = randint(1, N-1)
    g = 1
    while g==1:
        x = ((x*x)%N+c)%N
        y = ((y*y)%N+c)%N
        y = ((y*y)%N+c)%N
        g = gcd(abs(x-y),N)
    return g
```

This returns just a single divisor. To get all the divisors we can use a simple generator

```
def generator(N):
    while not isprime(N) and N > 1:
        factor = pollard_rho(N)
        for fac in pollard_rho_generator(factor):
            yield fac
        N //= factor
    yield N
```

We will use this generator for a few other functions as well. For now it is enough to say that it recursively generates the prime factors. A quick explanation of how this algorithm works can be found in the footnotes⁴

Brent's factorization method is an improvement to Pollard's rho algorithm, published by R. Brent in 1980 [9]. In Pollard's rho algorithm, one tries to find a non trivial factor s of N by finding indices i, j with $i < j$ such that $x_i \equiv x_j \pmod{s}$ and $x_i \not\equiv x_j \pmod{N}$. The x_n sequence is defined by the recurrence relation:

$$x_0 \equiv 2 \pmod{N}$$

$$x_{n+1} \equiv x_n^2 + 2 \pmod{N}$$

Pollard suggested that x_n be compared to x_{2n} for $n = 1, 2, 3, \dots$. Brent's improvement to Pollard's method is to compare x_n to x_m , where m is the largest integral power of 2 less than n .

Brent's factorization method

```
def brent(num):
    if num % 2 == 0:
        return 2
    y, c, m = randint(1, num-1), randint(1, num-1), randint(1, num-1)
    s, r, q = 1, 1, 1
    while s == 1:
        x = y
        for i in range(r):
            y = ((y*y) % num + c) % num
        k = 0
        while (k < r and s == 1):
```

³See https://en.wikipedia.org/wiki/Pollard%27s_rho_algorithm

⁴The Birthday Paradox: A Quick Tutorial on Pollard's Rho Algorithm <https://www.cs.colorado.edu/~srirams/courses/csci2824-spr14/pollardsRho.html>

```

        ys = y
        for i in range(min(m, r-k)):
            y = ((y*y) % num+c) % num
            q = q*(abs(x-y)) % num
            s = gcd(q, num)
            k += m
        r = r*2
    if s == num:
        while True:
            ys = ((ys*ys) % num+c) % num
            s = gcd(abs(x-ys), num)
            if s > 1:
                break
    return s

```

We can make a slight improvement to this algorithm. Most fast prime factorization algorithms start with sieving out the small prime factors.

```

from primesieve import generate_n_primes

NUM_OF_PRIMES = 10**3
PRIMES = generate_n_primes(NUM_OF_PRIMES)

def small_factor_generator(num):
    if isprime(num):
        yield num
    for prime in PRIMES:
        if num % prime == 0:
            num //= prime
            yield prime # Returns the next factor
            while num % prime == 0:
                num //= prime
                yield prime
            if prime > num**0.5:
                break
    yield num

```

As usual Python has specialized algorithms for factoring integers, one of these is the `primefac` package⁵. The description of the package reads as follows:

This is a module and command-line utility for factoring integers. As a module, we provide a primality test, several functions for extracting a non-trivial factor of an integer, and a generator that yields all of a number's prime factors (with multiplicity). As a command-line utility, this project aims to replace GNU's `factor` command with a more versatile utility — in particular, this utility can operate on arbitrarily large numbers, uses multiple cores in parallel, uses better algorithms, handles input in reverse Polish notation, and can be tweaked via command-line flags. Specifically

- One thread runs Brent's variation on Pollard's rho algorithm. This is good for extracting smallish factors quickly.
- One thread runs the two-stage version of Pollard's p-1 method. This is good at finding factors p for which p-1 is a product of small primes.
- One thread runs Williams' p+1 method. This is good at finding factors p for which p+1 is a product of small primes.

⁵<https://pypi.python.org/pypi/primefac>

- One thread runs the elliptic curve method. This is a bit slower than Pollard's rho algorithm when the factors extracted are small, but it has significantly better performance on difficult factors.
- One thread runs the multiple-polynomial quadratic sieve. This is the best algorithm for factoring "hard" numbers short of the horrifically complex general number field sieve. However, it's (relatively speaking) more than a little slow when the numbers are small, and the time it takes depends only on the size of the number being factored rather than the size of the factors being extracted as with Pollard's rho algorithm and the elliptic curve method, so we use the preceding algorithms to handle those.

There has been much studies done on using elliptic curves for factorizing integers. These methods in general are quite complex to not only implement but also understand. I implemented a naive version of the Lenstra elliptic curve factorization [?], however this did not give any speed increases on the integers tested. The reason is that the elliptic curve I used was choosen at random. The `primefac` package probably has a much better implementation than me.

Some of the heavier methods are much slower for small factor. The following list gives a rough estimate for which algorithm should be used at what number range.

- Small Numbers : Use simple sieve algorithms to create list of primes and do plain factorization. Works blazingly fast for small numbers.
- Big Numbers : Use Pollard's rho algorithm, Shanks' square forms factorization (Thanks to Dana Jacobsen for the pointer)
- Less Than 10^{25} : Use Lenstra elliptic curve factorization [?]
- Less Than 10^{100} : Use Use Quadratic sieve
- More Than 10^{100} : See [?] for a layman introduction to the General number field sieve [GNFS]. For a more in depth study of GNFS see [?].

A speed comparison for the more advanced factoring methods can be found in ???. Brent* denotes, uses a prime sieve to remove small primefactors before invoking Brent's improved Pollard Rho algorithm.

Table 2: Timings in ms for some of the prime factor methods

Number	Prime gen	Pollard Rho	Brent	Brent*
600 851 475 143	0,268	0,303	0,242	0,082
6 008 514 751 436 008	0,246	0,244	0,216	0,232
600 851 475 143 600 851 475 143	1,639	1,955	1,040	0,367
1 389 133 318 189	34,951	3,589	1,427	1,724
138 912 436 076 543	284,297	11,040	3,784	3,894
13 891 248 322 099 591	timeout	103,697	20,750	14,773

References

- [] Matthew E. Briggs, *An introduction to the general number field sieve*, Tech. report, 1998, http://www.math.vt.edu/people/brown/doc/briggs_gnfs_thesis.pdf.
- [] H. W. Lenstra, *Factoring integers with elliptic curves*, Annals of Mathematics **126** (1987), no. 3, 649–673, <http://www.jstor.org/stable/1971363>.
- [] Carl Pomerance, *A tale of two sieves*, Notices American Mathematical Society **43** (1996), 1473–1485, <http://www.ams.org/notices/199612/pomerance.pdf>.

Project Euler 4: Largest palindrome product

A palindromic number reads the same both ways. The largest palindrome made from the product of two 2-digit numbers is $9009 = 91 \times 99$.

Find the largest palindrome made from the product of two 3-digit numbers.

```
def palindrome_product_naive(length):
    palindrome_max = 0
    lower = 10**(length-1) + 1
    higher = 10**length
    for k in range(lower, higher):
        for j in range(lower, higher):
            if ispalindrome(str(j*k)):
                palindrome = j*k
                if palindrome > biggest_palindrome:
                    palindrome_max = palindrome
    print biggest_palindrome
```

51994.8730455 ms 445.112778322 ms

Project Euler 5: Smallest multiple

2520 is the smallest number that can be divided by each of the numbers from 1 to 10 without any remainder.

What is the smallest positive number that is evenly divisible by all of the numbers from 1 to 20?

some tekst

```
def smallest_multiple(number):
    prod = 1
    lst = [i for i in range(2,number+1)]
    for divisor in lst:
        if divisor > 1:
            prod *= divisor
            for k, num in enumerate(lst):
                if (num % divisor) == 0:
                    lst[k] /= divisor
    return prod
```

Some text

```
from primesieve import generate_primes
from math import log

def smallest_multiple_fast(n):
    prod = 1
    for prime in generate_primes(n):
        prod *= prime**int(log(n, prime))
    return prod
```

Project Euler 7: 10 001st prime

By listing the first six prime numbers: 2, 3, 5, 7, 11, and 13, we can see that the 6th prime is 13.

What is the 10 001st prime number?

Project Euler 8: Largest prime factor

The prime factors of 13195 are 5, 7, 13 and 29.

What is the largest prime factor of the number 600851475143?

Project Euler 3: Largest prime factor

The prime factors of 13195 are 5, 7, 13 and 29.

What is the largest prime factor of the number 600851475143?

Project Euler 10: Summation of primes

The sum of the primes below 10 is $2 + 3 + 5 + 7 = 17$.

Find the sum of all the primes below two million.

Project Euler 13: Large Sum

Work out the first ten digits of the sum of the following one-hundred 50-digit numbers.

```
37107287533902102798797998220837590246510135740250
46376937677490009712648124896970078050417018260538
74324986199524741059474233309513058123726617309629
91942213363574161575252240563301811072406154908250
23067588207539346171171980310421047513778063246676
89261670696623633820136378418383684178734361726757
28112879812849979408065481931592621691275889832738
4427422891743252032192358942287696487670272189318
47451445736001306439091167216856844588711603153276
70386486105843025439939619828917593665686757934951
62176457141856560629502157223196586755079324193331
64906352462741904929101432445813822663347944758178
92575867718337217661963751590579239728245598838407
58203565325359399008402633568948830189458628227828
80181199384826282014278194139940567587151170094390
35398664372827112653829987240784473053190104293586
8651550606295864861532075273371959191420517255829
71693888707715466499115593487603532921714970056938
5437007057682668462462149565007647178729443837604
53282654108756828443191190634694037855217779295145
36123272525000296071075082563815656710885258350721
45876576172410976447339110607218265236877223636045
17423706905851860660448207621209813287860733969412
81142660418086830619328460811191061556940512689692
51934325451728388641918047049293215058642563049483
62467221648435076201727918039944693004732956340691
15732444386908125794514089057706229429197107928209
55037687525678773091862540744969844508330393682126
18336384825330154686196124348767681297534375946515
8038628759287849020152168554828717201219257766954
78182833757993103614740356856449095527097864797581
16726320100436897842553539920931837441497806860984
48403098129077791799088218795327364475675590848030
87086987551392711854517078544161852424320693150332
59959406895756536782107074926966537676326235447210
69793950679652694742597709739166693763042633987085
41052684708299085211399427365734116182760315001271
65378607361501080857009149939512557028198746004375
35829035317434717326932123578154982629742552737307
94953759765105305946966067683156574377167401875275
88902802571733229619176668713819931811048770190271
25267680276078003013678680992525463401061632866526
36270218540497705585629946580636237993140746255962
24074486908231174977792365466257246923322810917141
91430288197103288597806669760892938638285025333403
34413065578016127815921815005561868836468420090470
23053081172816430487623791969842487255036638784583
11487696932154902810424020138335124462181441773470
6378329949063625966498587618221225225512486764533
67720186971698544312419572409913959008952310058822
95548255300263520781532296796249481641953868218774
76085327132285723110424803456124867697064507995236
37774242535411291684276865538926205024910326572967
23701913275725675285653248258265463092207058596522
29798860272258331913126375147341994889534765745501
18495701454879288984856827726077713721403798879715
38298203783031473527721580348144513491373226651381
34829543829199918180278916522431027392251122869539
40957953066405232632538044100059654939159879593635
29746152185502371307642255121183693803580388584903
41698116222072977186158236678424689157993532961922
6246795719440126904387710725048102390895523597457
23189706772547915061505504953922979530901129967519
8618808822587531452958409925120382900940777075672
11306739708304724483816533873502340845647058077308
82959174767140363198008187129011875491310547126581
97623331044818386269515456334926366572897563400500
42846280183517070527831839425882145521227251250327
55121603546981200581762165212827652751691296897789
323819573432933994643750190780936333018982642090
75506164965184775180738168837861091527357929701337
62177842752192623401942399639168044983993173312731
32924185707147349566916674687634660915035914677504
99518671430235219628894890102423325116913619626622
73267460800591547471830798392868535206946944540724
7684182252467441716151403642798227334805556214818
97142617910342598647204516893989422179826088076852
8778364618279934631376775430780936333018982642090
10848802521674670883215120185883543223812876952786
71329612474782464538636993009049310363619763878039
62184073572399794223406235393808339651327408011116
66627891981488087797941876876144230030984490851411
60661826293682836764744779239180335110989069790714
85786944089552990653640447425576083659976645795096
66024396409905389607120198219976047599490197230297
64913982680032973156037120041377903785566085089252
16730939319872750275468906903707539413042652315011
94809377245048795150954100921645863754710598436791
78639167021187492431995700641917969777599028300699
15368713711936614952811305876380278410754449733078
40789923115535562561142322423255033685442488917353
44889911501440648020369068063960672322193204149535
41503128880339536053299340368006977710650566631954
81234880673210146739058568557934581403627822703280
82616570773948327592232845941706525094512325230608
22918802058777319719839450180888072429661980811197
77158542502016545090413245809786882778948721859617
72107838435069186155435662884062257473692284509516
20849603980134001723930671666823555245252804609722
53503534226472524250874054075591789781264330331690
```

While this problem is quite easy; it introduces a handful of ideas and techniques which will be very useful for handling harder problems. All the solutions below assumes that the numbers have been saved in a textfile, with each number on a new line. With this in mind a simple solution reads

```
def sum_first_n_digits(filename, digits = 10):
    return str(sum(int(line) for line in open(filename, 'r')))[0:digits]
```

This reads the entire textfile into a list, before taking the sum of that list. This sum is converted into a string, and the first digits are returned. One problem with the above code is that it reads in the entire file

before summing each number. This can consume a lot of memory what if we had several billion numbers? With this in mind a solution is to read the file line adding each number to a running counter

```
def sum_first_n_digits_2(filename, digits = 10):
    total = 0
    with open(filename, "r") as file_of_numbers:
        for number in file_of_numbers:
            total += int(number)
    return str(total)[0:digits]
```

Some might wonder why I do not explicitly close the file after opening it. This is because the file in the example above actually gets with `open(filename, "r") as file_of_numbers:` is a context processor - this means that Python takes care of freeing the resources, according to the "context manager" protocol, which file object adheres to. See preshing.com/20110920/the-python-with-statement-by-example and python.org/dev/peps/pep-0343 to learn more.

String conversion is as we will see inherently slow and something which is desirable to avoid. However in this cases the alternatives are just as slow. With some clever use of modular arithmetic we can get the first n -digits of a number as

```
def first_n_digits(number, digits):
    total_digits = int(math.log(number, 10)+1)
    insignificant = 10**(total_digits-digits)
    return (number - number%insignificant)/insignificant
```

Another optimization one might try is to reduce the number of arithmetical operations that is performed. Since we are only interested in the first 10-digits, why not just sum the first 11 digits of each number? Indeed, for this case of numbers it works, however no noticable speed improvement is gained. This is again due to string operations being slow.

```
def sum_first_n_digits_3(filename, digits = 10):
    total = 0
    with open(filename, "r") as file_of_numbers:
        for number in file_of_numbers:
            total += int(number[0:digits+1])
    return str(total)[0:digits]
```

However this solution does not always produce the correct results. Take a minute to think why we can not simply add the first 11 numbers, to get the first 10 right. A counterexample is shown below:

```
999123
999438
999439
```

Assume we want to figure out the first 4 digits of the sum of these three numbers. Let us try to sum the first 4, 5 and 6 digits of each number

4	:	9991	+	9994	+	9994	=	29979
5	:	99912	+	99943	+	99943	=	299798
6	:	999123	+	999438	+	999439	=	2998000

As the math below show the only way to avoid getting round off errors is to sum every number in its entirety.